

C4Scale Takehome Coding Assignment

Technical Interview Assignment

Overview

Create an AI-powered question-answering agent that can process documentation from help websites and accurately answer user queries about product features, integrations, and functionality.

Requirements Core Functionality

- The agent should accept a help website URL as input (e.g., help.zluri.com, help.slack.com)
- Process and index the content for efficient querying
- Accept natural language questions via terminal interface
- Provide accurate answers based on the processed documentation
- Clearly indicate when information is not available in the documentation

Technical Requirements

Input Processing

- Handle different help website URL formats
- Implement proper error handling for invalid URLs
- Support recursive crawling of documentation pages

Content Processing

- Extract meaningful content while filtering out navigation elements, headers, footers
- Maintain proper context hierarchy of the documentation
- Handle different content types (text, lists, tables)

Question Answering

- Process natural language questions
- Implement semantic search or similar technology for finding relevant content
- Format answers in a clear, readable manner
- Include source references (URLs) for answers when possible

Error Handling

- Graceful handling of network issues
- Clear feedback for unsupported websites
- Proper indication when questions cannot be answered

Example Usage

Starting the agent

```
python qa_agent.py --url https://help.example.com
```

Example interaction

> What integrations are available?

[Agent responds with integration information from documentation]

> How do I enable feature X?

[Agent provides step-by-step instructions if available]

> Does the product support feature Y?

Evaluation Criteria 1. Accuracy (50%)

- Correctness of answers
- Relevance of responses
- Proper handling of "no information" cases
- Consistency in answers
- Context preservation

2. Technical Implementation and Innovation (35%)

- Efficient approach
- Appropriate choice of search/matching algorithm
- Resource usage optimization
- Response time
- Scalability considerations

3. Code Quality (15%)

- Clean, well-organized code structure
- Proper error handling and edge cases
 - Sorry, I couldn't find any information about feature Y in the documentation.
- Use of appropriate design patterns
- Code documentation and comments
- Modular and maintainable architecture
- Proper dependency management

Submission Requirements

1. Code Repository

- Submit via a private GitHub repository
- Include clear README with:
 - Setup instructions
 - Dependencies
 - Usage examples
 - Design decisions
 - Known limitations

2. Documentation

- Technical architecture overview
- Implementation approach
- Future improvement suggestions
- Testing approach

3. Testing

- Include unit tests
- Provide test cases with example questions and expected answers
- Include performance benchmarks

Time Expectation

- 48 hours for a basic implementation

Bonus Points

- 1. Advanced Features
 - Support for multiple documentation sources
 - Answer caching mechanism
 - Support for different documentation formats
 - Confidence scores for answers
- 2. Technical Improvements
 - Docker containerization
 - API endpoint addition
 - Performance optimizations
 - Advanced NLP techniques

Notes

- You may use any programming language, but Python is preferred
- You can use any LLM APIs, RAG Approaches, any vector DBs, Agentic framework of your choice
- Document any third-party libraries used
- Include any assumptions made during implementation
- Note any limitations of your approach

Submission Process

1. Share repository access with cvp@c4scale.com
2. Include a brief video demo (5 minutes max)
3. Be prepared to discuss your implementation decisions in the review call